


# PATTERN RECOGNITION USING PYTHON

## Classification



Wen-Yen Hsu

Dept Electrical Engineering  
Chang Gung University, Taiwan

2019-Spring

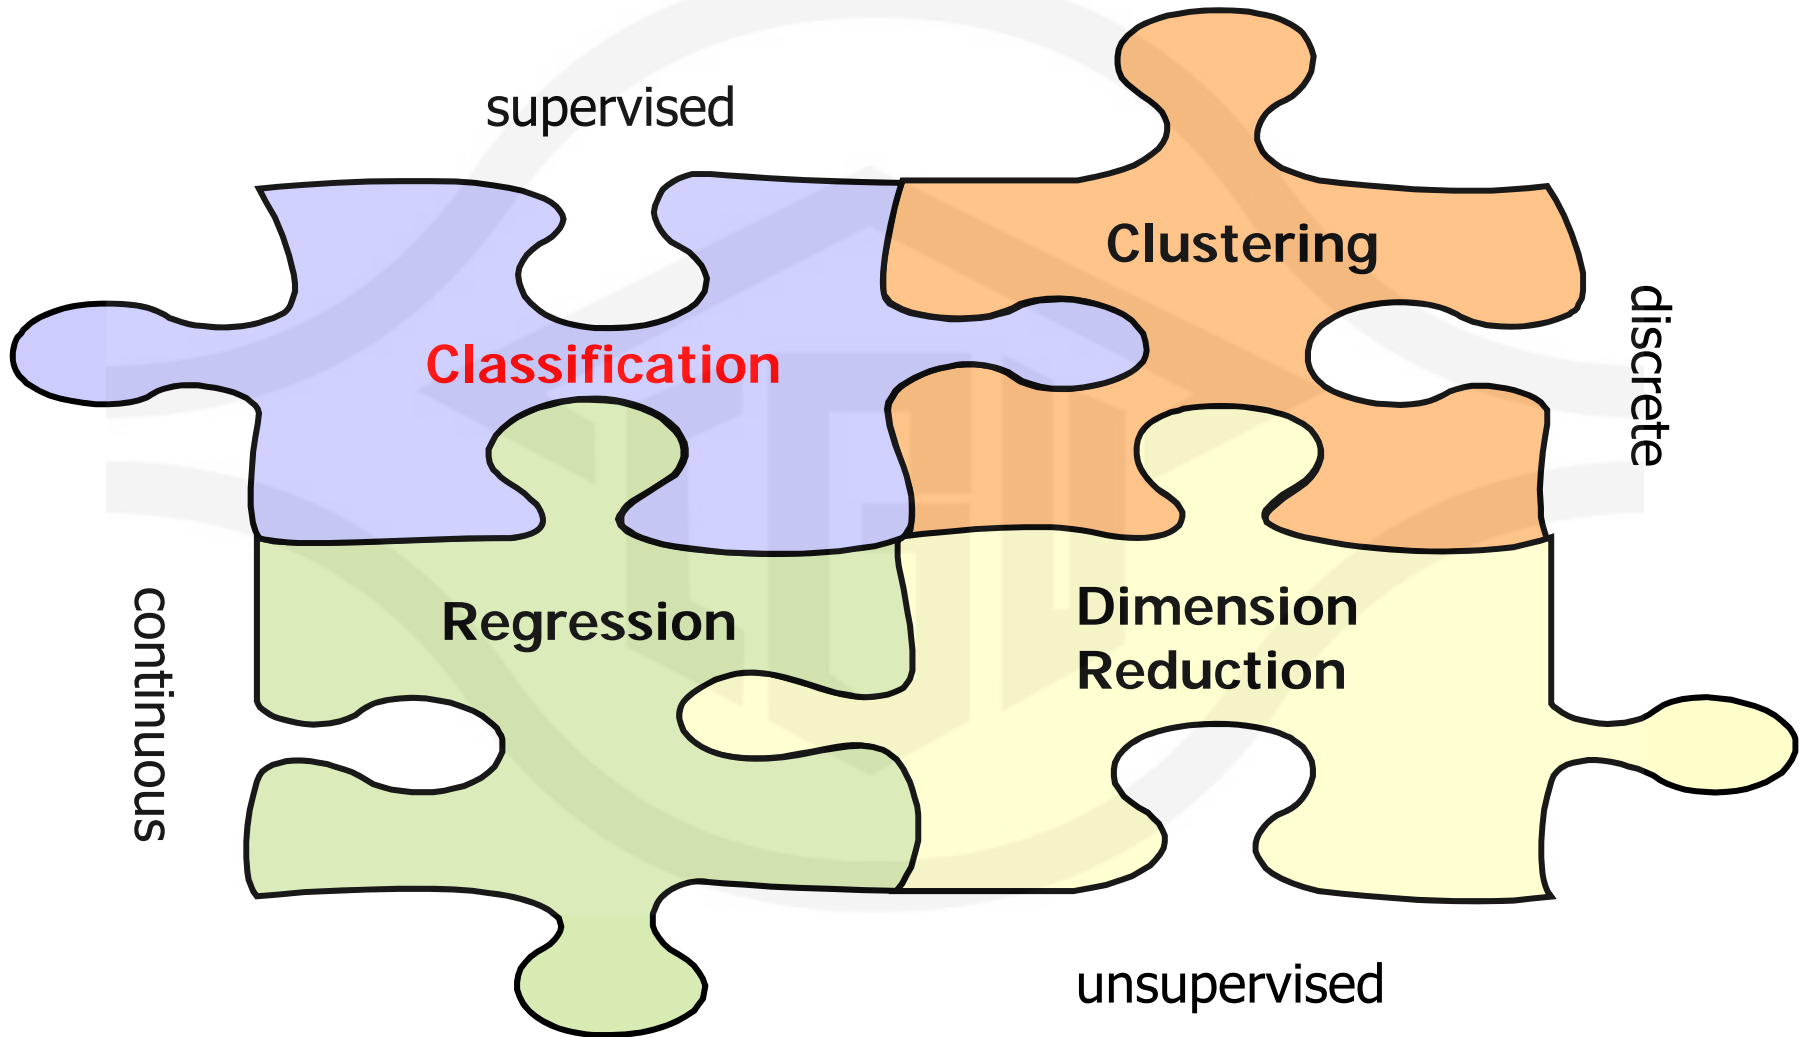
# Supervised Learning Algorithms for Classification

---

- Popular algorithms for classification, such as logistic regression, support vector machines, and decision trees
- Examples and explanations using the scikit-learn machine learning library, which provides a wide variety of machine learning algorithms via a user-friendly Python API

# Machine Learning Organizational Chart

---

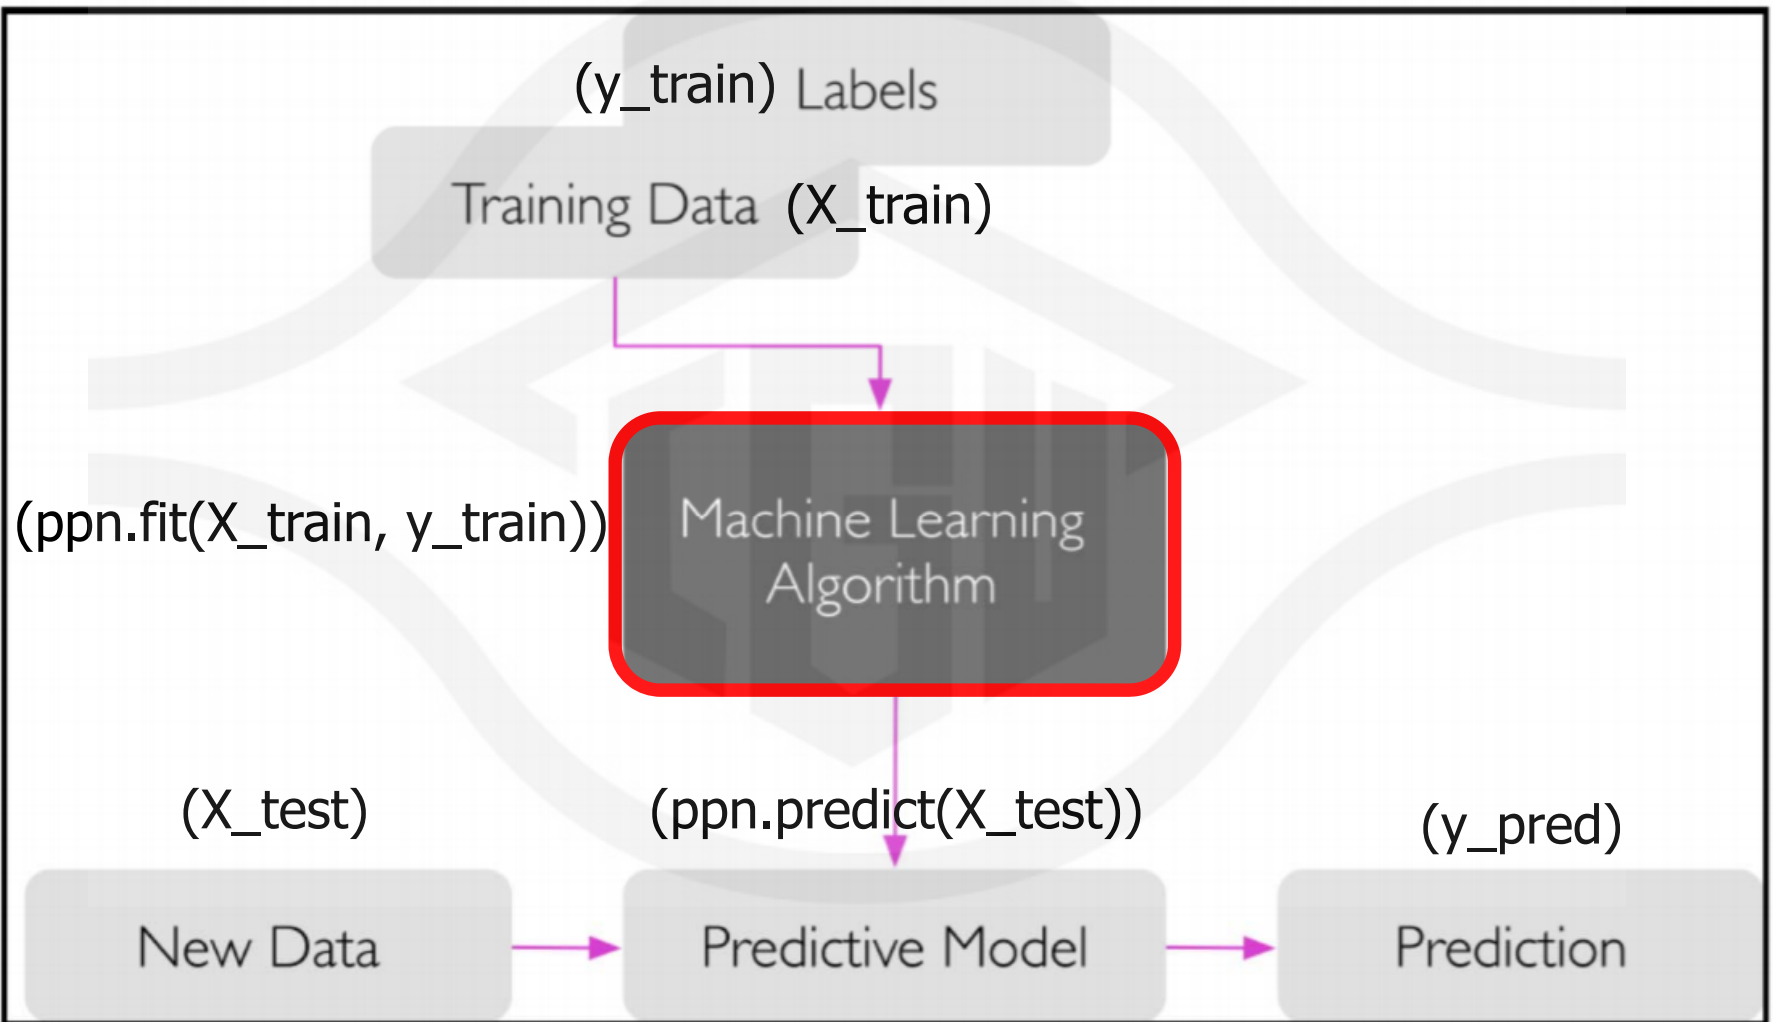


# Choosing a Classification Algorithm

---

- Selecting features and collecting training samples
- Choosing a performance metric
- Choosing a classifier and optimization algorithm
- Evaluating the performance of the model
- Tuning the algorithm

# Solve Classification Problem



# Problem Transformation

## Machine Learning Algorithm

Classification Problem



evaluate **error** (when label  $\neq$  prediction)  
by cost function (objective function)

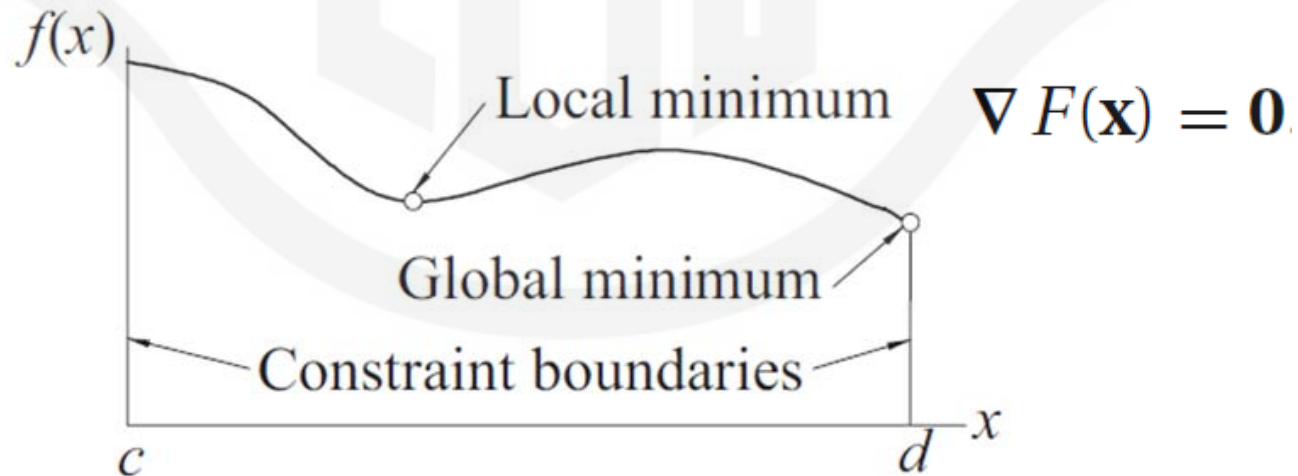
Optimization Problem

# Optimization Problem

- In engineering, optimization is closely related to design.
- We wish to keep it **(objective function) as small as possible**, such as the cost or weight.

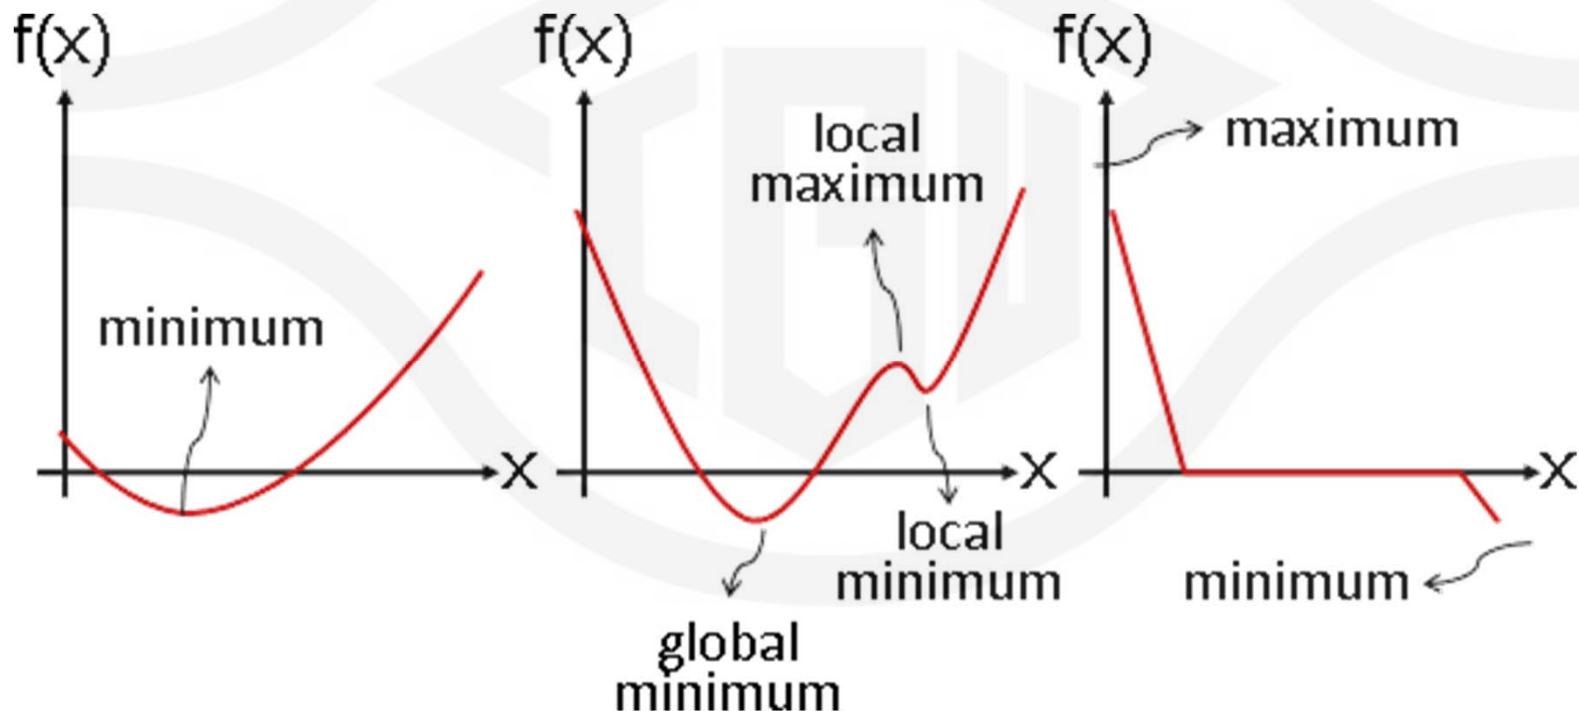
Find  $\mathbf{x}$  that minimizes  $F(\mathbf{x})$  subject to  $g(\mathbf{x}) = 0$ ,  $h(\mathbf{x}) \geq 0$ .

- Find the points where the gradient vector of  $F(\mathbf{x})$  vanishes



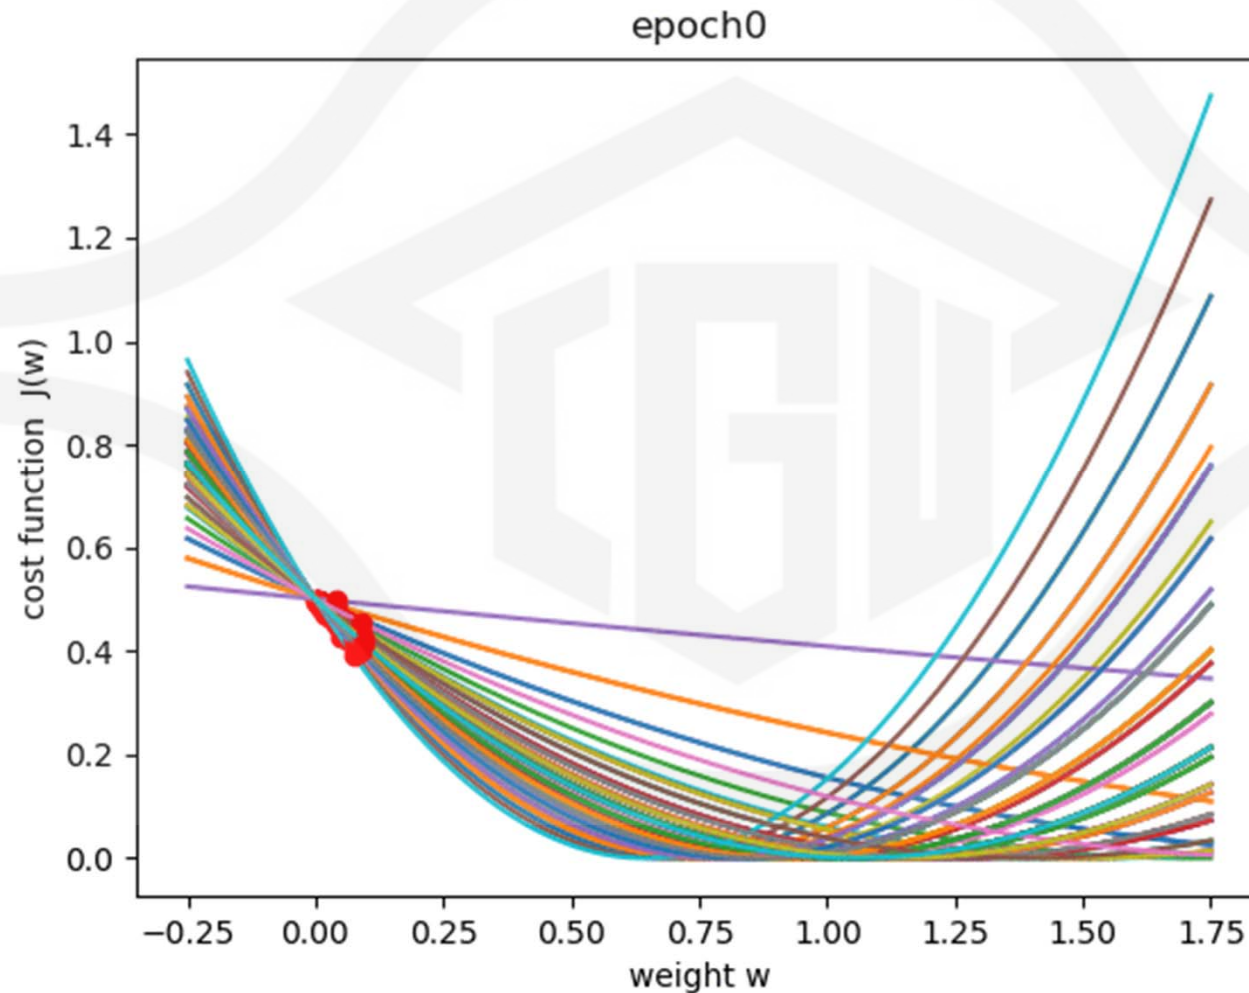
# Algorithms for Minimization

- It require **starting values** (`self.w_`)
- Processes by **iterative procedures** (`n_iter`)







# Gradient Descent to Achieve Minimization



# Effect of Learning Rate

- Gradient descent with different learning rate

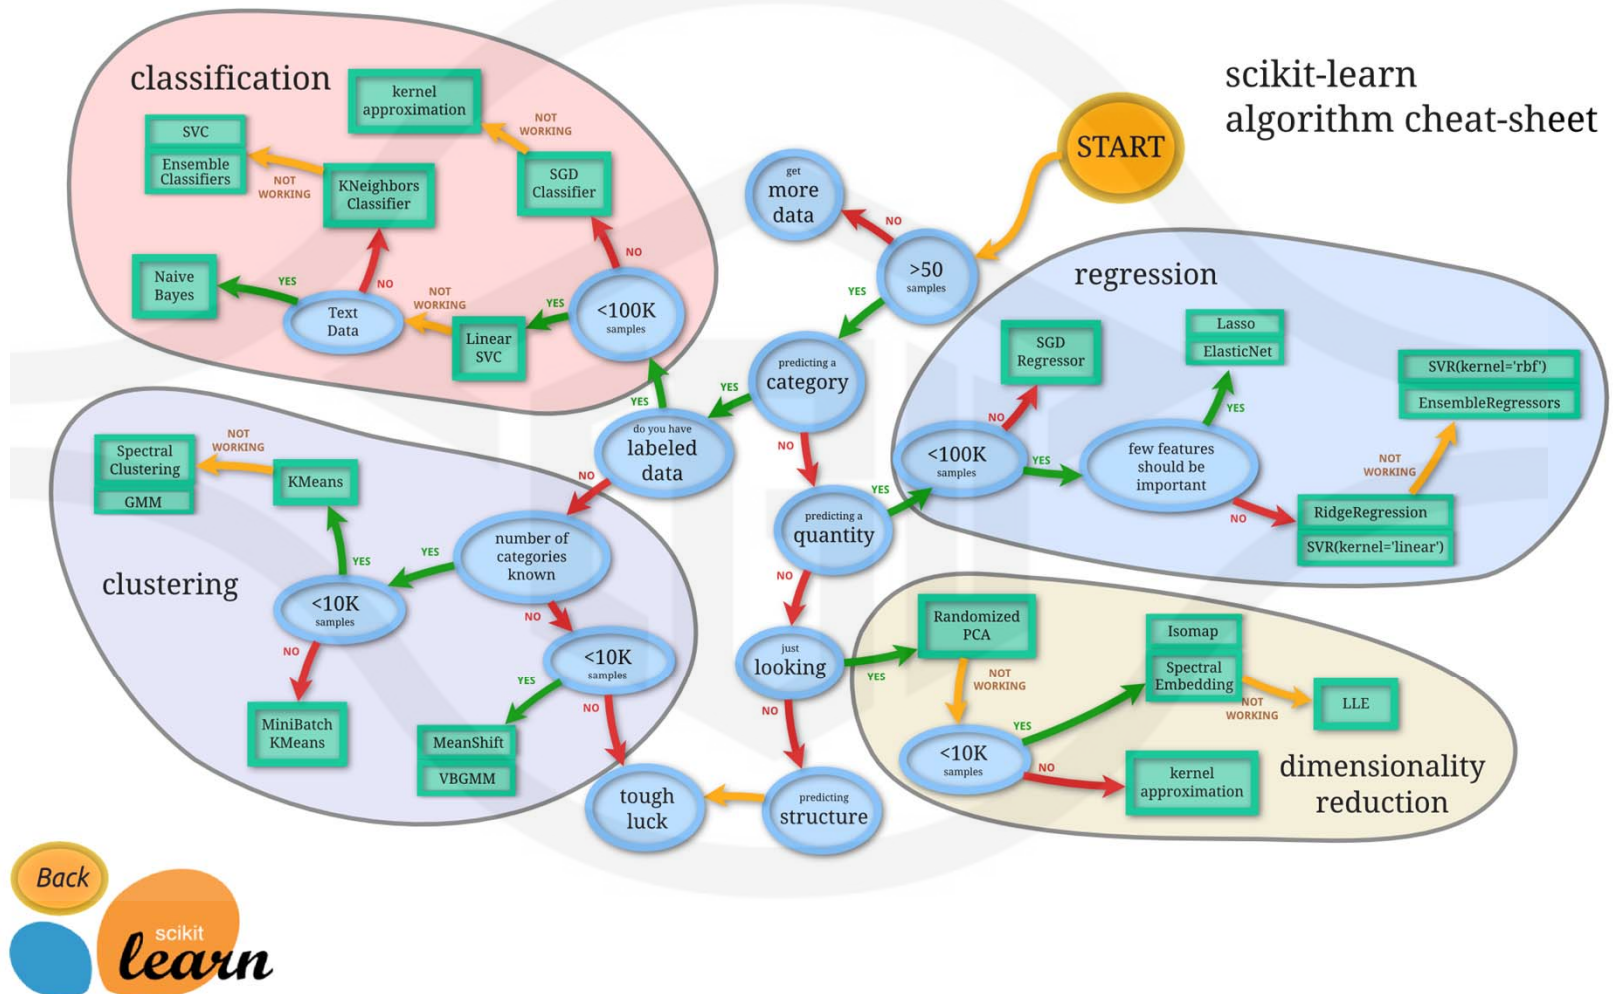


# What is Scikit-Learn

- <https://scikit-learn.org/stable/>



# Learning Map of Scikit-Learn



# Import Scikit-Learn (sklearn)

---

- from ... import ... (special use method)

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score
```

# Load Data via sklearn

- Load data by built-in scikit-learn datasets

```
iris = datasets.load_iris()  
X = iris.data[:, [2, 3]]  
y = iris.target  
print('Class labels:', np.unique(y))
```

The iris dataset is a classic and very easy multi-class classification dataset.

Classes	3
Samples per class	50
Samples total	150
Dimensionality	4
Features	real, positive

# Preprocessing via sklearn

---

- Split data (70% train set and 30% test set)

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.3, random_state=1, stratify=y)
```

- Standardize features

```
sc = StandardScaler()  
sc.fit(X_train)  
X_train_std = sc.transform(X_train)  
X_test_std = sc.transform(X_test)
```

# Training and Testing

---

- Initialize the object and training

```
ppn = Perceptron(n_iter=40, eta0=0.1, random_state=1)
ppn.fit(X_train_std, y_train)
```

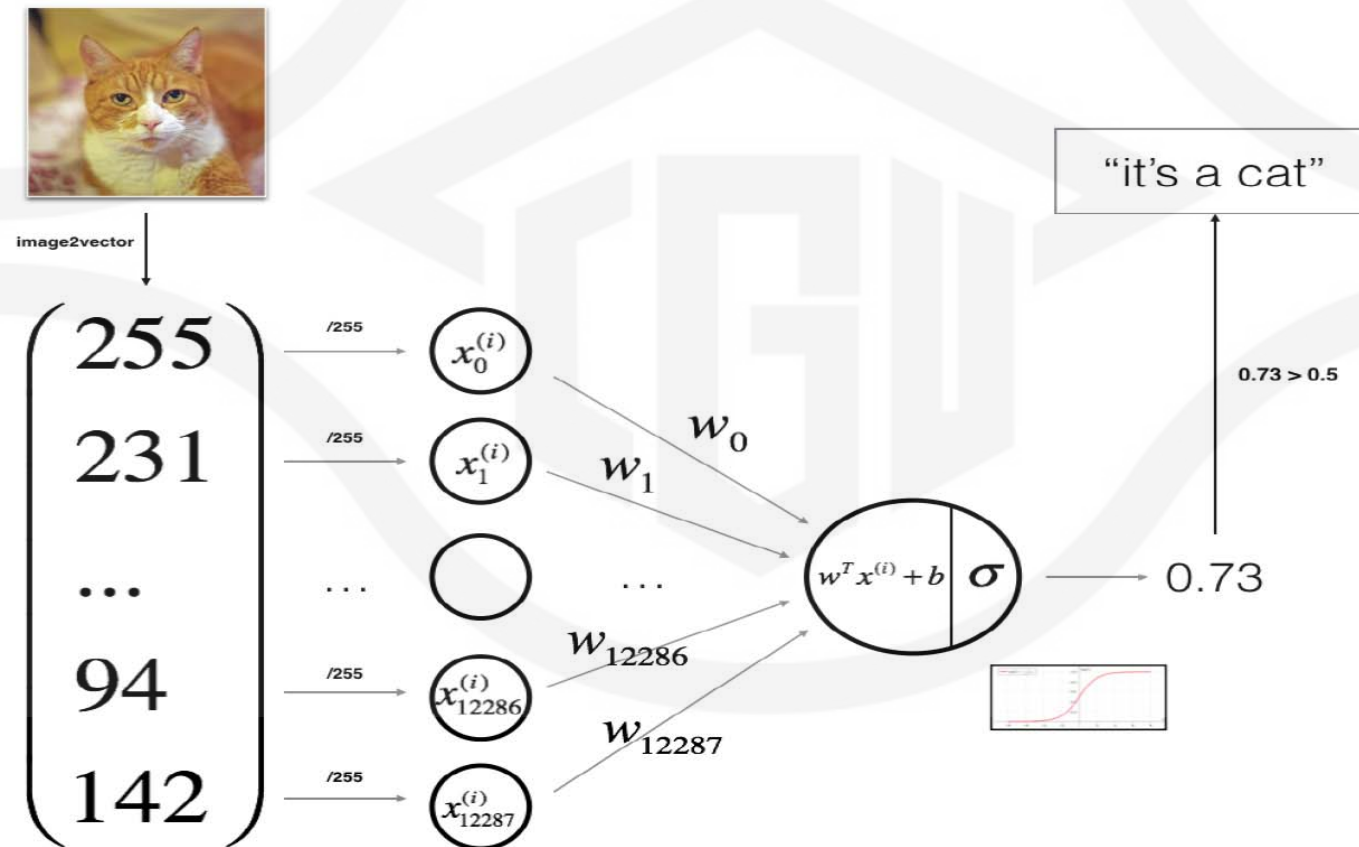
- Check the accuracy of test set

```
y_pred = ppn.predict(X_test_std)
print('Misclassified samples: %d' % (y_test !=
y_pred).sum())
print('Accuracy: %.2f' % accuracy_score(y_test, y_pred))
```

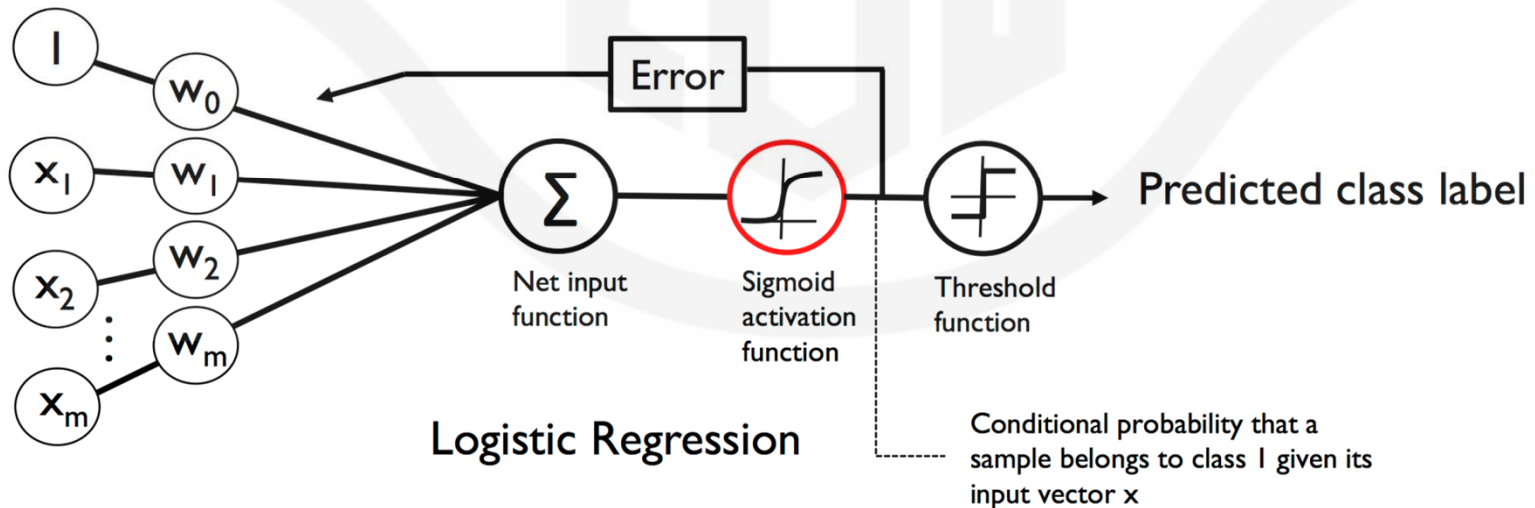
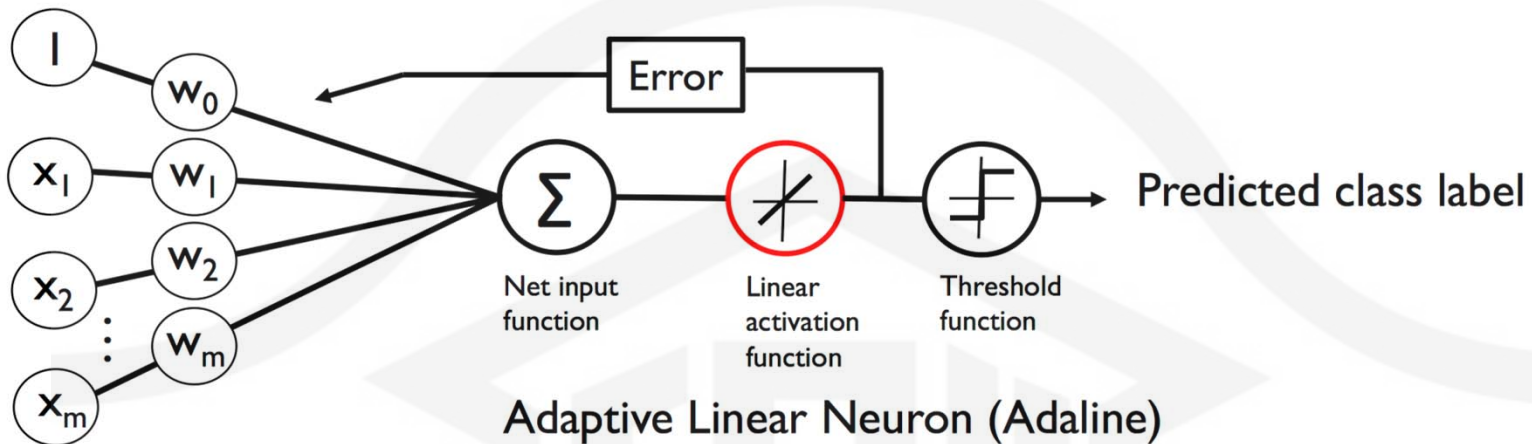


# Logistic Regression


- Logistic Regression is actually a very simple Neural Network



# Compare with Adaline



# Activation Function with Probability



- Link  $\phi(z)$  to Bernoulli's PDF  $P(y | x, w)$

$$P(y|x, w) = p^y (1 - p)^{1-y} \quad \begin{cases} p, \text{ if } y = 1 \\ 1 - p, \text{ if } y = 0 \end{cases}$$

# Logistic Regression Characteristic

---

- Replace with the sigmoid function

```
def activation(self, z):  
    """Compute logistic sigmoid activation"""  
    return 1. / (1. + np.exp(-z))
```

- Prediction range constrained between 0 to 1

```
def predict(self, X):  
    """Return class label after unit step"""  
    return np.where(self.activation(self.net_input(X))  
    >= 0.5, 1, 0)
```

# Maximum Likelihood Estimation

- If training examples were identically independently distributed (IID), then maximum likelihood estimation (MLE) means to find the parameters of the model with given training data


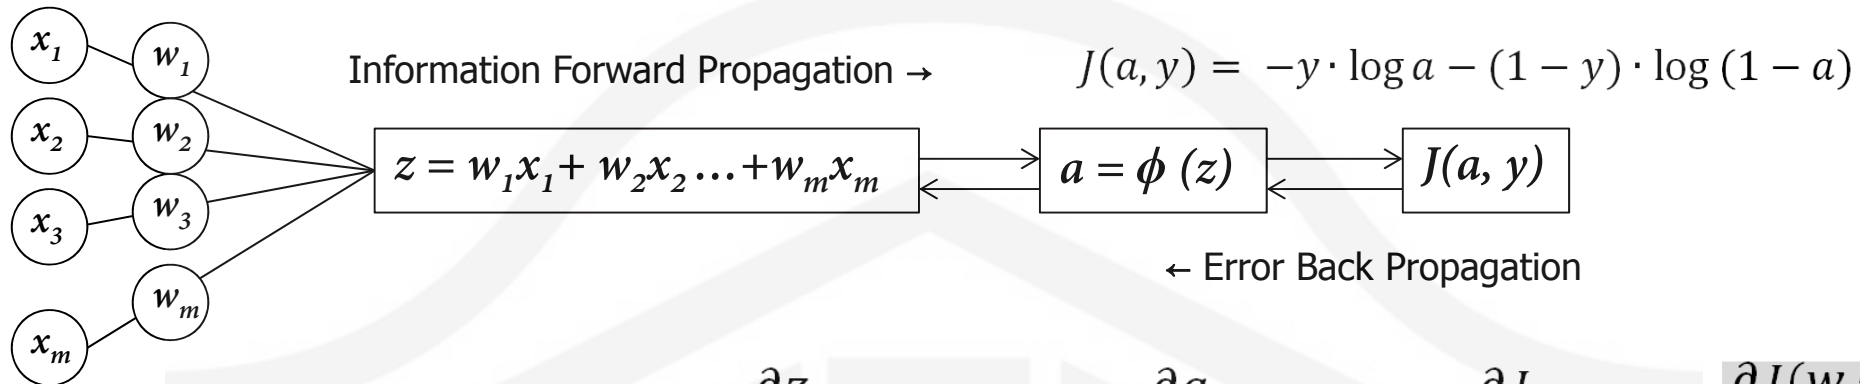
$$L(\mathbf{w}) = P(\mathbf{y} | \mathbf{x}; \mathbf{w}) = \prod_{i=1}^n P(y^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}) = \prod_{i=1}^n \left( \phi(z^{(i)}) \right)^{y^{(i)}} \left( 1 - \phi(z^{(i)}) \right)^{1-y^{(i)}}$$

$$l(\mathbf{w}) = \log L(\mathbf{w}) = \sum_{i=1}^n \left[ y^{(i)} \log \left( \phi(z^{(i)}) \right) + (1 - y^{(i)}) \log \left( 1 - \phi(z^{(i)}) \right) \right]$$

- Negative log-likelihood as the cost function

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^{(i)} \log \left( \phi(z^{(i)}) \right) - (1 - y^{(i)}) \log \left( 1 - \phi(z^{(i)}) \right) \right]$$

# Learning Phase



$$\frac{\partial z}{\partial w_j} \times \frac{\partial a}{\partial z} \times \frac{\partial J}{\partial a} = \frac{\partial J(w_j)}{\partial w_j}$$

$$\downarrow \quad \quad \downarrow \quad \quad \downarrow$$

$$(x_j) \quad (a \cdot (1 - a)) \quad \left(-\frac{y}{a} + \frac{1-y}{1-a}\right)$$

$$\Rightarrow \frac{\partial J(w_j)}{\partial w_j} = (a - y)x_j = \nabla J(w_j)$$

$$w_j^{Next} = w_j^{Now} - \eta \cdot \nabla J(w_j)$$

# About Cost Function

- **Logistic regression or conditional log-likelihood cost function** ( $-\log P(y|x)$ ) worked much better than the **quadratic cost**(squared errors) which was traditionally used to train feedforward neural networks for classification problems.

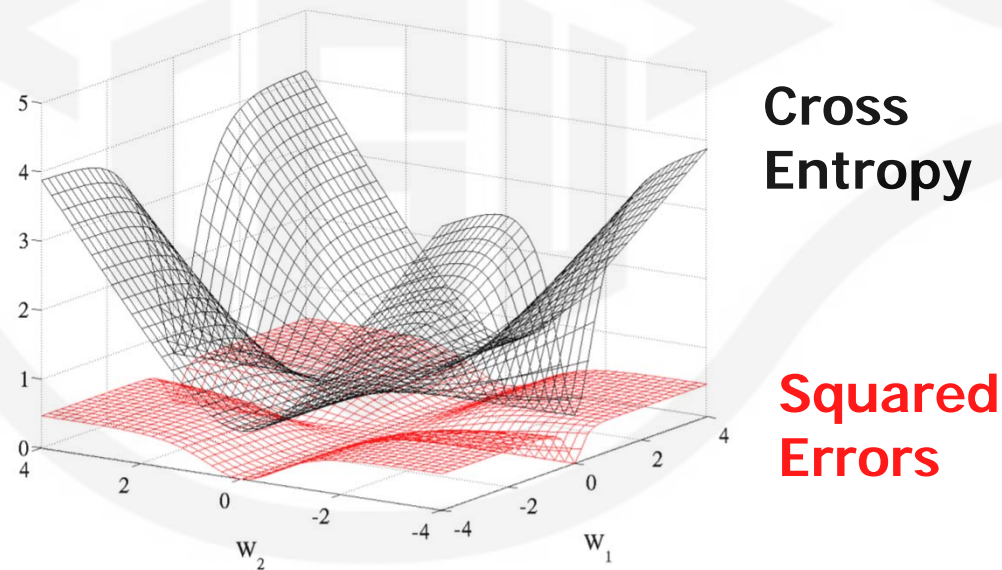


Figure 5: Cross entropy (black, surface on top) and quadratic (red, bottom surface) cost as a function of two

# Logistic Regression Training Algorithm

## ■ Implement

```
def fit(self, X, y):
    rgen = np.random.RandomState(self.random_state)
    self.w_ = rgen.normal(loc=0.0, scale=0.01, size=1 + X.shape[1])
    self.cost_ = []
    for _ in range(self.n_iter):
        net_input = self.net_input(X)
        output = self.activation(net_input)
        errors = (y - output)
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()
        # note that we compute the logistic `cost` now
        # instead of the sum of squared errors cost
        cost = -y.dot(np.log(output)) - ((1 - y).dot(np.log(1 - output)))
        self.cost_.append(cost)
    return self
```




# Logistic Regression via sklearn

- Initialize the object and training

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(C=100.0, random_state=1)
lr.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined,
                      classifier=lr, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

# Tackling Overfitting via Regularization



- L2 regularization


$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right] + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

# Regularization Parameter

- $C$  is regularization parameter  $\lambda$ 's inverse

```
weights, params = [], []
for c in np.arange(-5, 5):
    lr = LogisticRegression(C=10.**c, random_state=1)
    lr.fit(X_train_std, y_train)
    weights.append(lr.coef_[2])
    params.append(10.**c)
weights = np.array(weights)
plt.plot(params, weights[:, 0],
         label='petal length')
plt.plot(params, weights[:, 1], linestyle='--',
         label='petal width')
plt.ylabel('weight coefficient')
plt.xlabel('C')
plt.legend(loc='upper left')
plt.xscale('log')
plt.show()
```

# Support Vector Machine



$$w_0 + \mathbf{w}^T \mathbf{x}_{pos} = 1 \quad (1)$$

$$w_0 + \mathbf{w}^T \mathbf{x}_{neg} = -1 \quad (2)$$

$$\Rightarrow \mathbf{w}^T (\mathbf{x}_{pos} - \mathbf{x}_{neg}) = 2$$

$$\|\mathbf{w}\| = \sqrt{\sum_{j=1}^m w_j^2}$$

$$\frac{\mathbf{w}^T (\mathbf{x}_{pos} - \mathbf{x}_{neg})}{\|\mathbf{w}\|} = \frac{2}{\|\mathbf{w}\|}$$

$$w_0 + \mathbf{w}^T \mathbf{x}^{(i)} \geq 1 \text{ if } y^{(i)} = 1$$

$$w_0 + \mathbf{w}^T \mathbf{x}^{(i)} \leq -1 \text{ if } y^{(i)} = -1$$

$$\text{for } i = 1 \dots N$$



# Linear Support Vector Machine via sklearn

- Initialize the object and training

```
from sklearn.svm import SVC
svm = SVC(kernel='linear', C=1.0, random_state=1)
svm = SVC(kernel='rbf', random_state=1, gamma=0.10,
C=10.0)
svm.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std,
                      y_combined,
                      classifier=svm,
                      test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

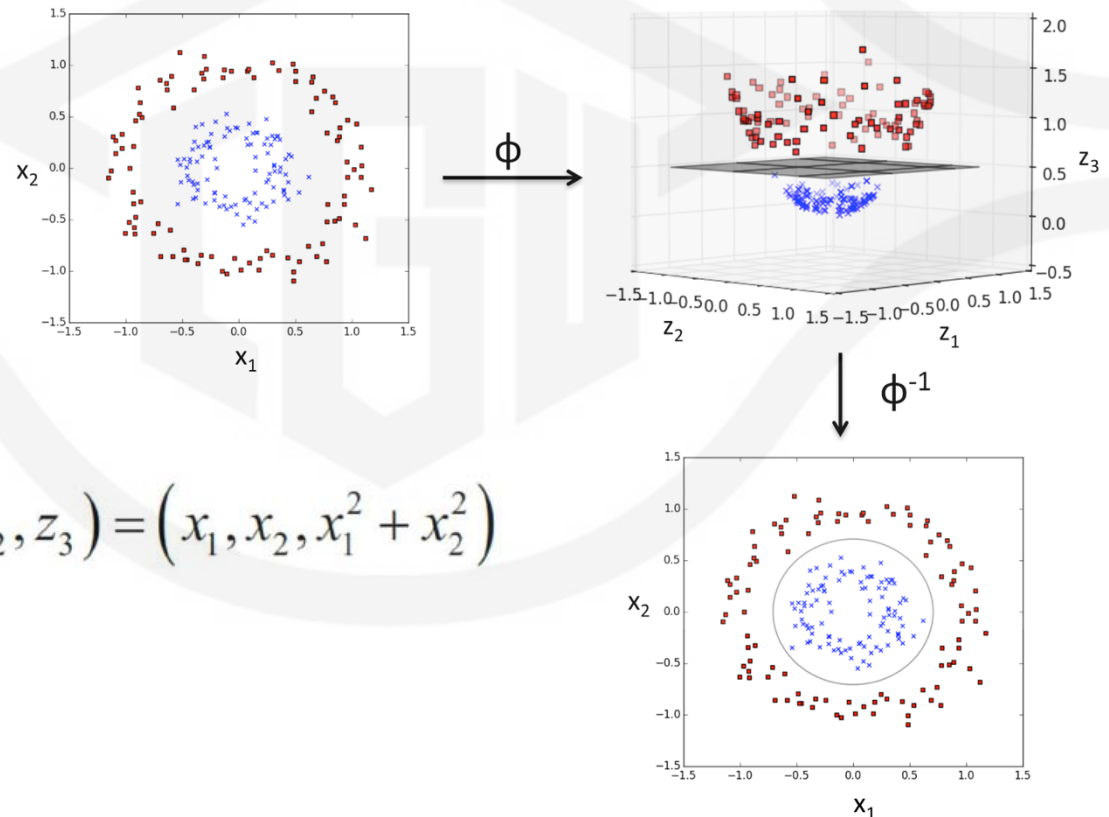
# Non-Linear Problems

- Exclusive-OR problems



# Kernel Methods

- Create nonlinear combinations of original features to project them onto a higher-dimensional space where it becomes linearly separable.



$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

# Separating Hyperplanes in High-dimensional Space

---

- Increase features' dimension is computationally expensive especially if dealing with high-dimensional data
- Kernel trick
  - Replace the dot product  $\mathbf{x}^{(i)T} \mathbf{x}^{(j)}$  by  $K(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$
- Radial Basis Function (RBF) kernel or Gaussian kernel
  - Similarity function between a pair of samples.
  - A range between 1 (for exactly similar samples) and 0 (for very dissimilar samples)

$$\mathcal{K}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\gamma \left\| \mathbf{x}^{(i)} - \mathbf{x}^{(j)} \right\|^2\right)$$



# Using the Kernel Trick

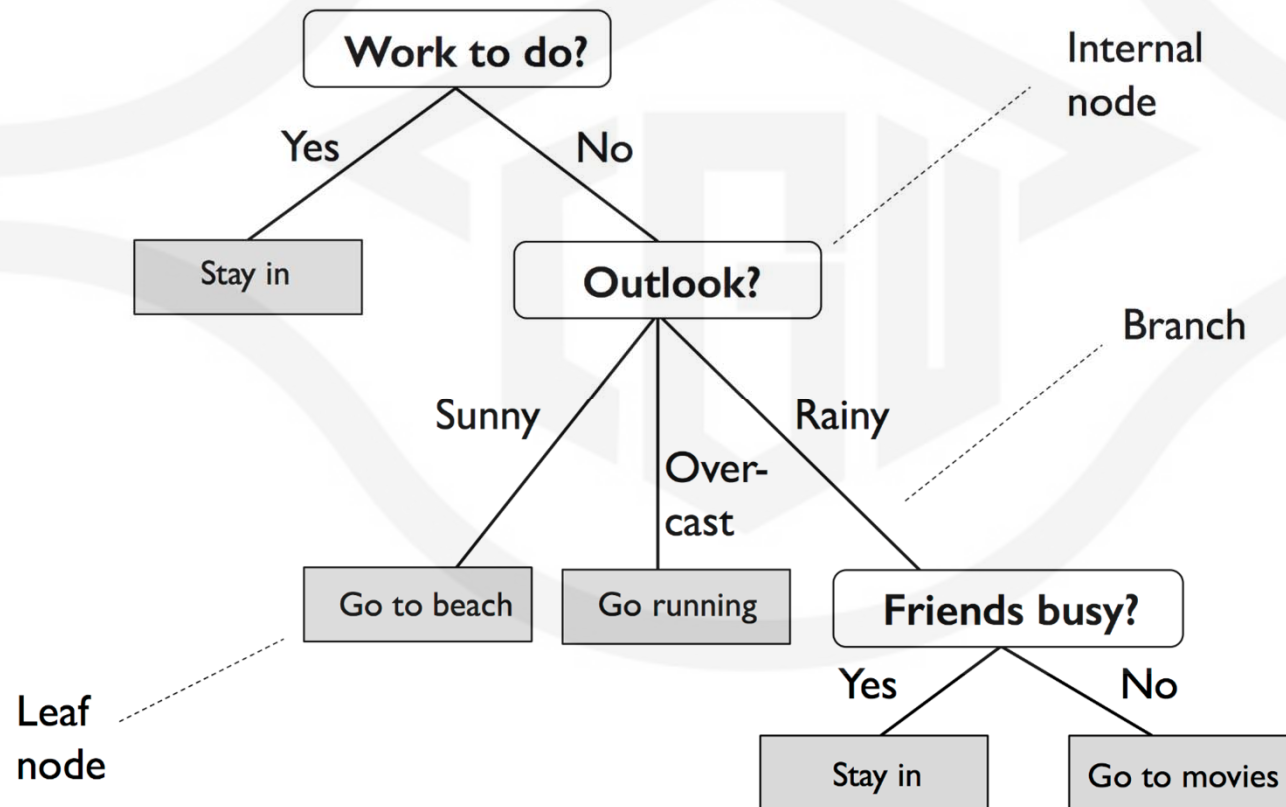
## ■ RBF kernel

```
np.random.seed(1)
X_xor = np.random.randn(200, 2)
y_xor = np.logical_xor(X_xor[:, 0] > 0,
                        X_xor[:, 1] > 0)
y_xor = np.where(y_xor, 1, -1)
svm = SVC(kernel='rbf', random_state=1, gamma=0.10,
           C=10.0)
svm.fit(X_xor, y_xor)
plot_decision_regions(X_xor, y_xor,
                      classifier=svm)

plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

# Decision Tree

- **Interpretability:** make a decision based on asking a series of questions



# Information Gain

---

- Using the decision algorithm, we start at the tree root and split the data on the feature that results in the largest **Information Gain (IG)**
- **Maximizing information gain**

$$IG(D_p, f) = I(D_p) - \frac{N_{left}}{N_p} I(D_{left}) - \frac{N_{right}}{N_p} I(D_{right})$$

# Impurity Measure by Entropy (IH)

---

$$I_H(t) = -\sum_{i=1}^c p(i|t) \log_2 p(i|t)$$

- $p(i|t)$ : proportion of samples that belong to class  $i$  for a node  $t$ 
  - The entropy is 0 if all samples belong to the same class
  - The entropy is maximal if a uniform class distribution
  - For example, entropy=0 if  $p(i=1|t)=1, p(i=0|t)=0$ ; entropy=1 if  $p(i=1|t)=0.5, p(i=0|t)=0.5$ .
- **The entropy criterion attempts to maximize the mutual information in the tree**

# Gini Impurity (IG)

---

$$I_G(t) = \sum_{i=1}^c p(i|t)(1 - p(i|t)) = 1 - \sum_{i=1}^c p(i|t)^2$$

- Gini impurity is maximal if the classes are perfectly mixed, for example, in a binary class setting (  $c = 2$  ):

$$I_G(t) = 1 - \sum_{i=1}^c 0.5^2 = 0.5$$

# Impurity Measure by Classification Error (IE)

---

$$I_E = 1 - \max \{ p(i | t) \}$$

- A useful criterion for pruning but not recommended for growing a decision tree, since it is less sensitive to changes in the class probabilities of the nodes.

# Decision Tree via sklearn

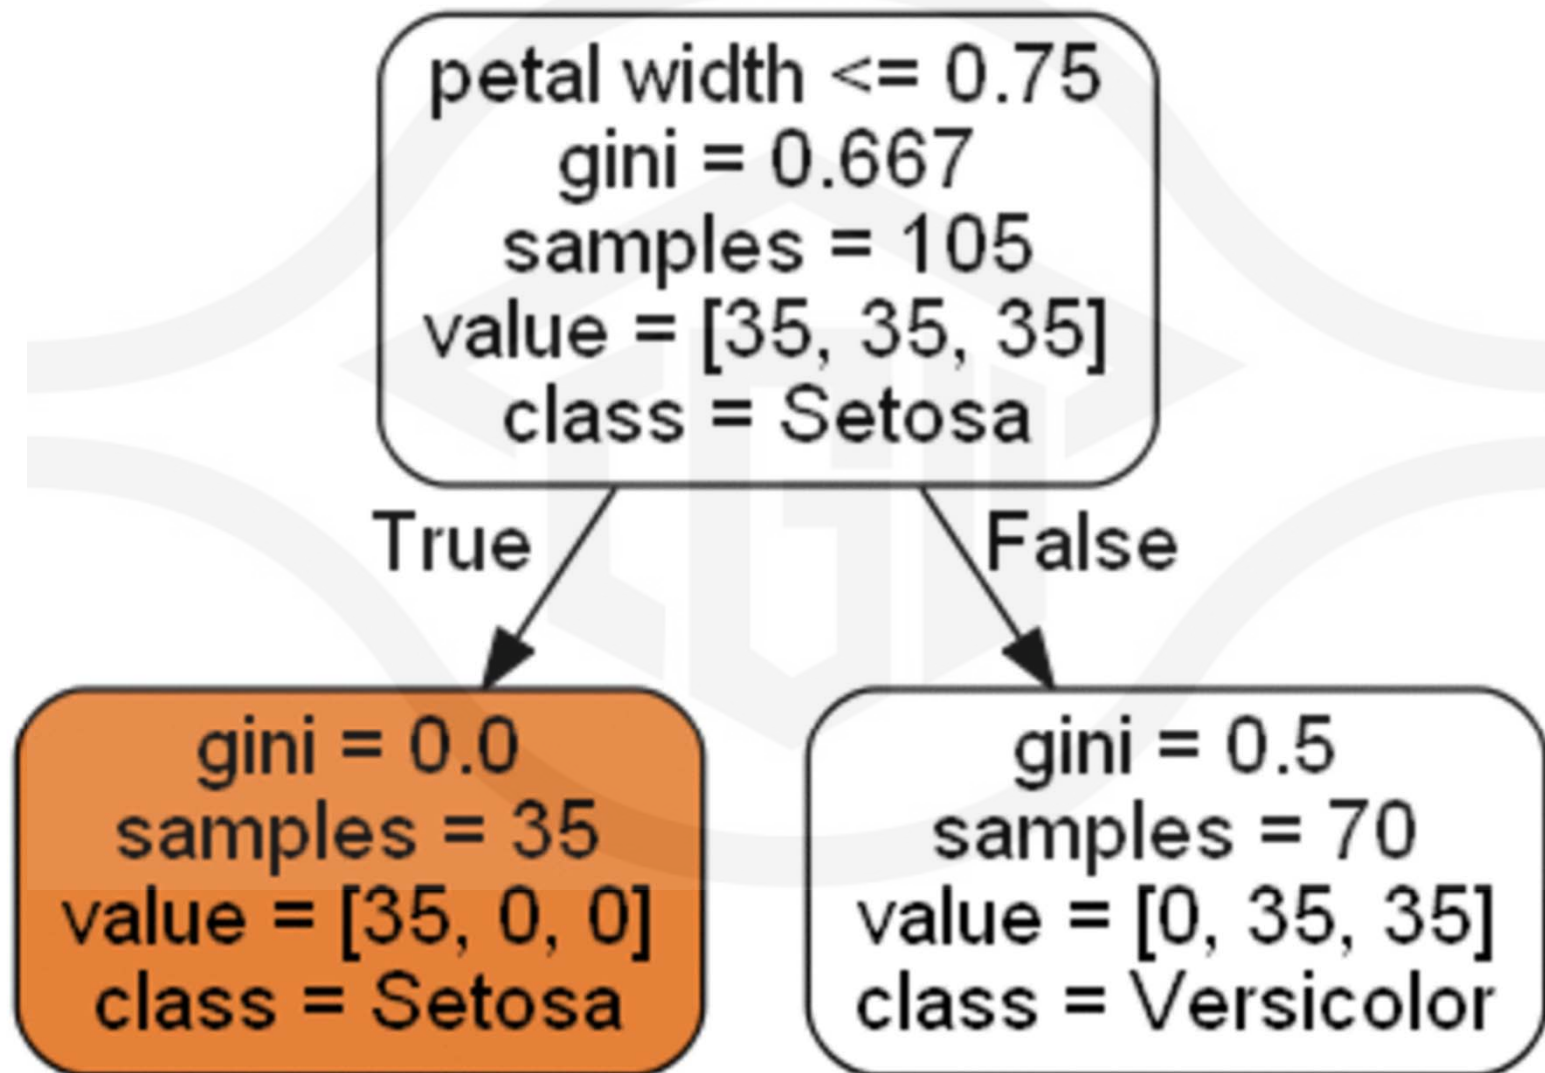
---

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(criterion='gini',
                             max_depth=4,
                             random_state=1)

tree.fit(X_train, y_train)
X_combined = np.vstack((X_train, X_test))
y_combined = np.hstack((y_train, y_test))
plot_decision_regions(X_combined, y_combined,
                     classifier=tree,
                     test_idx=range(105, 150))
plt.xlabel('petal length [cm]')
plt.ylabel('petal width [cm]')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

# Max Depth

---





# Multiple Decision Trees to Random Forests

---

- Random forest algorithm
  - Draw a random **bootstrap** sample (randomly choose  $n$  samples from the training set).
  - Grow a decision tree from the bootstrap sample. At each node:
    - Randomly select  $d$  features.
    - Split the node using the feature that provides the best split (maximal information gain).
  - Repeat the steps 1-2  $k$  times.
  - Aggregate the prediction by each tree to assign the class label by **majority vote**.

# Random Forests via sklearn


---

```
from sklearn.ensemble import RandomForestClassifier
forest = RandomForestClassifier(criterion='gini',
                               n_estimators=25,
                               random_state=1,
                               n_jobs=2)

forest.fit(X_train, y_train)
plot_decision_regions(X_combined, y_combined,
                      classifier=forest,
                      test_idx=range(105, 150))
plt.xlabel('petal length [cm]')
plt.ylabel('petal width [cm]')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

# K-Nearest Neighbors

- It doesn't learn a discriminative function from the training data, but memorizes the training dataset instead
  - Choose the number of  $k$  and a distance metric.
  - Find the  $k$ -nearest neighbors of the sample to classify.
  - Assign the class label by majority vote.



# Minkowski distance

---

- Minkowski distance is typically used with  $p$  being 1 or 2, which correspond to the Manhattan distance and the Euclidean distance, respectively.

$$D(X, Y) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

# K-Nearest Neighbors via sklearn

---

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5,
                          p=2,
                          metric='minkowski')
knn.fit(X_train_std, y_train)
plot_decision_regions(X_combined_std, y_combined,
                      classifier=knn,
                      test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

# Reference

---

- Sebastian Raschka, Vahid Mirjalili. Python Machine Learning: Machine Learning and Deep Learning with Python, scikit-learn, and TensorFlow. Second Edition. Packt Publishing, 2017.